

# CS 595: Assignment #8

Due on Thursday, November 13, 2014

*Dr. Nelson 4:20pm*

Holly Harkins

## Contents

Problem 1	3
-----------	---

## Problem 1

The MovieLens data sets were collected by the GroupLens Research Project at the University of Minnesota during the seven-month period from September 19th, 1997 through April 22nd, 1998. It is available for download from <http://www.grouplens.org/node/73>

There are three files which we will use:

1. `u.data`: 100,000 ratings by 943 users on 1,682 movies. Each user has rated at least 20 movies. Users and items are numbered consecutively from 1. The data is randomly ordered. This is a tab separated list of

```
user id | item id | rating | timestamp
```

The time stamps are unix seconds since 1/1/1970 UTC.

2. `u.item`: Information about the 1,682 movies. This is a tab separated list of

```
movie id | movie title | release date | video release date | IMDb URL  
| unknown | Action | Adventure | Animation | Children's | Comedy | Crime  
| Documentary | Drama | Fantasy | Film-Noir | Horror | Musical | Mystery  
| Romance | Sci-Fi | Thriller | War | Western |
```

The last 19 fields are the genres, a 1 indicates the movie is of that genre, a 0 indicates it is not; movies can be in several genres at once. The movie ids are the ones used in the `u.data` data set.

3. `u.user`: Demographic information about the users. This is a tab separated list of:

```
user id | age | gender | occupation | zip code
```

The user ids are the ones used in the `u.data` data set.

The code for reading from the `u.data` and `u.item` files and creating recommendations is described in the book *Programming Collective Intelligence* (check email for more details). You are to modify `recommendations.py` to answer the following questions. Each question your program answers correctly will award you 1 point.

Answers:

To do this assignment, I modified `recommendations.py`. I added the functionality for each question to the `loadMovieLens()` function. `Assignment8.py` uses the data from `u.data`, `u.item`, and `u.user` to output the results. I had some trouble with question 7. I do not believe 4 user1 agreed most with the same user 98. I also used the tabular view for the first time. My code highlights how I answered the questions. It is listed at the end of the document.

1. What 5 movies have the highest average ratings? Show the movies and their ratings sorted by their average ratings.

Title	Rating
Great Day in Harlem, A (1994)	5.0
Prefontaine (1997)	5.0
Aiqing wansui (1994)	5.0
Star Kid (1997)	5.0
Marlene Dietrich: Shadow and Light (1996)	5.0
Entertaining Angels: The Dorothy Day Story (1996)	5.0
Saint of Fort Washington, The (1993)	5.0
Someone Else's America (1995)	5.0
Santa with Muscles (1996)	5.0
They Made Me a Criminal (1939)	5.0

2. What 5 movies received the most ratings? Show the movies and the number of ratings sorted by number of ratings.

Title	Ratings Count
Star Wars (1977)	583
Contact (1997)	509
Fargo (1996)	508
Return of the Jedi (1983)	507
Liar Liar (1997)	485

3. What 5 movies were rated the highest on average by women? Show the movies and their ratings sorted by ratings.

Title	Rating
Visitors, The (Visiteurs, Les) (1993)	5.0
Prefontaine (1997)	5.0
Telling Lies in America (1997)	5.0
Foreign Correspondent (1940)	5.0
Faster Pussycat! Kill! Kill! (1965)	5.0
Year of the Horse (1997)	5.0
Mina Tannenbaum (1994)	5.0
Maya Lin: A Strong Clear Vision (1994)	5.0
Everest (1998)	5.0
Someone Else's America (1995)	5.0
Stripes (1981)	5.0

4. What 5 movies were rated the highest on average by men? Show the movies and their ratings sorted by ratings.

Title	Rating
Delta of Venus (1994)	5.0
Great Day in Harlem, A (1994)	5.0
Leading Man, The (1996)	5.0
Love Serenade (1996)	5.0
Prefontaine (1997)	5.0
Aiqing wansui (1994)	5.0
Little City (1998)	5.0
Star Kid (1997)	5.0
Marlene Dietrich: Shadow and Light (1996)	5.0
Entertaining Angels: The Dorothy Day Story (1996)	5.0
Quiet Room, The (1996)	5.0
Saint of Fort Washington, The (1993)	5.0
Letter From Death Row, A (1998)	5.0
Santa with Muscles (1996)	5.0
They Made Me a Criminal (1939)	5.0

5. What movie received ratings most like Top Gun? Which movie received ratings that were least like Top Gun (negative correlation)?

Most Like Top Gun : 1.0, 'Wild America (1997)'

Least Like Top Gun: -1.0, 'Babysitter, The (1995)'

6. Which 5 raters rated the most films? Show the raters' IDs and the number of films each rated..

Rater ID	Ratings Count
405	736
655	678
13	632
450	538
276	516

7. Which 5 raters most agreed with each other? Show the raters' IDs and Pearson's  $r$ , sorted by  $r$ .

User1 ID	User2 ID	Pearson's $r$
675	99	1.0
156	98	1.0
772	98	1.0
226	98	1.0
227	98	1.0

8. Which 5 raters most disagreed with each other (negative correlation)? Show the raters' IDs and Pearson's  $r$ , sorted by  $r$ .

User1 ID	User2 ID	$r$ Value
655	384	0.683
13	46	0.688
130	511	0.725
327	816	0.772
796	205	0.791

9. a. What movie was rated highest on average by men over 40?

Title	Rating
Delta of Venus (1994)	5.0
Santa with Muscles (1996)	5.0
Crossfire (1947)	5.0
Leading Man, The (1996)	5.0
Love Serenade (1996)	5.0
Prefontaine (1997)	5.0
Aiqing wansui (1994)	5.0
Love in the Afternoon (1957)	5.0
Star Kid (1997)	5.0
Angel Baby (1995)	5.0
Maya Lin: A Strong Clear Vision (1994)	5.0
Entertaining Angels: The Dorothy Day Story (1996)	5.0
Magic Hour, The (1998)	5.0
Quiet Room, The (1996)	5.0
Saint of Fort Washington, The (1993)	5.0
Perfect Candidate, A (1996)	5.0
Letter From Death Row, A (1998)	5.0
Little Princess, The (1939)	4.5
Grosse Fatigue (1994)	4.5
Sum of Us, The (1994)	4.5
Boy's Life 2 (1997)	4.5
Fille seule, La (A Single Girl) (1995)	4.5
Winter Guest, The (1997)	4.5
Man of No Importance, A (1994)	4.5
Anna (1996)	4.5
Two or Three Things I Know About Her (1966)	4.5
Innocents, The (1961)	4.5
Wallace and Gromit Aardman Animation (1996)	4.5
Casablanca (1942)	4.5
Paths of Glory (1957)	4.5

9. b. By men under 40?

Title	Rating
Hearts and Minds (1996)	5.0
Faithful (1996)	5.0
Marlene Dietrich: Shadow and Light (1996)	5.0
Strawberry and Chocolate (Fresa y chocolate) (1993)	5.0
Late Bloomers (1996)	5.0
Solo (1996)	5.0
Grateful Dead (1995)	5.0
Prefontaine (1997)	5.0
Rendezvous in Paris (Rendez-vous de Paris, Les) (1995)	5.0
World of Apu, The (Apu Sansar) (1959)	5.0
Aparajito (1956)	5.0
Ace Ventura: When Nature Calls (1995)	5.0
Star Kid (1997)	5.0
Two or Three Things I Know About Her (1966)	5.0
Poison Ivy II (1995)	5.0
Double Happiness (1994)	5.0
Little City (1998)	5.0
Boxing Helena (1993)	5.0
Spice World (1997)	5.0
They Made Me a Criminal (1939)	5.0
Great Day in Harlem, A (1994)	5.0
Little Princess, The (1939)	5.0
Unstrung Heroes (1995)	5.0
Leading Man, The (1996)	5.0
Indian Summer (1996)	5.0
Pather Panchali (1955)	4.8
A Chef in Love (1996)	4.7
Whole Wide World, The (1996)	4.7
Close Shave, A (1995)	4.7
Shanghai Triad (Yao a yao dao waipo qiao) (1995)	4.6

10. a. What movie was rated highest on average by women over 40? By women under 40?

Title	Rating
Shallow Grave (1994)	5.0
Great Dictator, The (1940)	5.0
Visitors, The (Visiteurs, Les) (1993)	5.0
Shall We Dance? (1937)	5.0
In the Bleak Midwinter (1995)	5.0
Funny Face (1957)	5.0
Ma vie en rose (My Life in Pink) (1997)	5.0
Swept from the Sea (1997)	5.0
Best Men (1997)	5.0
Foreign Correspondent (1940)	5.0
Tombstone (1993)	5.0
Wrong Trousers, The (1993)	5.0
Top Hat (1935)	5.0
Quest, The (1996)	5.0
Balto (1995)	5.0
Angel Baby (1995)	5.0
Band Wagon, The (1953)	5.0
Letter From Death Row, A (1998)	5.0
Mina Tannenbaum (1994)	5.0
Mary Shelley's Frankenstein (1994)	5.0
Gold Diggers: The Secret of Bear Mountain (1995)	5.0
Nightmare Before Christmas, The (1993)	5.0
Grand Day Out, A (1992)	5.0
Bride of Frankenstein (1935)	5.0
Pocahontas (1995)	5.0
Safe (1995)	5.0

10. b. By women under 40?

Title	Rating
Backbeat (1993)	5.0
Prefontaine (1997)	5.0
Telling Lies in America (1997)	5.0
Year of the Horse (1997)	5.0
Mina Tannenbaum (1994)	5.0
Maya Lin: A Strong Clear Vision (1994)	5.0
Nico Icon (1995)	5.0
Umbrellas of Cherbourg, The (Parapluies de Cherbourg, Les) (1964)	5.0
Everest (1998)	5.0
Heaven's Prisoners (1996)	5.0
Wedding Gift, The (1994)	5.0
Faster Pussycat! Kill! Kill! (1965)	5.0
Horseman on the Roof, The (Hussard sur le toit, Le) (1995)	5.0
Grace of My Heart (1996)	5.0
Someone Else's America (1995)	5.0
Don't Be a Menace to South Central While Drinking Your Juice in the Hood (1996)	5.0
Stripes (1981)	5.0



Listing 1: Modified recommendations.py

```
from __future__ import division
# A dictionary of movie critics and their ratings of a small
# set of movies
critics={
5 'Lisa Rose': {'Lady in the Water': 2.5, 'Snakes on a Plane': 3.5, 'Just My Luck':
    3.0, 'Superman Returns': 3.5, 'You, Me and Dupree': 2.5, 'The Night Listener':
    3.0},
'Gene Seymour': {'Lady in the Water': 3.0, 'Snakes on a Plane': 3.5, 'Just My Luck':
    1.5, 'Superman Returns': 5.0, 'The Night Listener': 3.0, 'You, Me and Dupree':
    3.5},
'Michael Phillips': {'Lady in the Water': 2.5, 'Snakes on a Plane': 3.0, 'Superman
    Returns': 3.5, 'The Night Listener': 4.0},
'Claudia Puig': {'Snakes on a Plane': 3.5, 'Just My Luck': 3.0, 'The Night Listener':
    4.5, 'Superman Returns': 4.0, 'You, Me and Dupree': 2.5},
'Mick LaSalle': {'Lady in the Water': 3.0, 'Snakes on a Plane': 4.0, 'Just My Luck':
    2.0, 'Superman Returns': 3.0, 'The Night Listener': 3.0, 'You, Me and Dupree': 2.0},
10 'Jack Matthews': {'Lady in the Water': 3.0, 'Snakes on a Plane': 4.0, 'The Night
    Listener': 3.0, 'Superman Returns': 5.0, 'You, Me and Dupree': 3.5},
'Toby': {'Snakes on a Plane': 4.5, 'You, Me and Dupree': 1.0, 'Superman Returns': 4.0}

from math import sqrt
15 import numpy as np
import operator
from collections import Counter

20 # Returns a distance-based similarity score for person1 and person2
def sim_distance(prefs, person1, person2):
    # Get the list of shared items
    si={}
    for item in prefs[person1]:
25         if item in prefs[person2]: si[item]=1

    # if they have no ratings in common, return 0
    if len(si)==0: return 0

30 # Add up the squares of all the differences
    sum_of_squares=sum([pow(prefs[person1][item]-prefs[person2][item],2)
                        for item in prefs[person1] if item in prefs[person2]])

    return 1/(1+sum_of_squares)

35 # Returns the Pearson correlation coefficient for p1 and p2
def sim_pearson(prefs, p1, p2):
    # Get the list of mutually rated items
    si={}
40     for item in prefs[p1]:
        if item in prefs[p2]: si[item]=1

    # if they are no ratings in common, return 0
    if len(si)==0: return 0
```

```
45     # Sum calculations
    n=len(si)

    # Sums of all the preferences
50    sum1=sum([prefs[p1][it] for it in si])
    sum2=sum([prefs[p2][it] for it in si])

    # Sums of the squares
    sum1Sq=sum([pow(prefs[p1][it],2) for it in si])
55    sum2Sq=sum([pow(prefs[p2][it],2) for it in si])

    # Sum of the products
    pSum=sum([prefs[p1][it]*prefs[p2][it] for it in si])

60    # Calculate r (Pearson score)
    num=pSum-(sum1*sum2/n)
    den=sqrt((sum1Sq-pow(sum1,2)/n)*(sum2Sq-pow(sum2,2)/n))
    if den==0: return 0

65    r=num/den

    return round(r,3)

# Returns the best matches for person from the prefs dictionary.
70 # Number of results and similarity function are optional params.
def topMatches(prefs, person, n=1682, similarity=sim_pearson):
    scores=[(similarity(prefs, person, other), other)
              for other in prefs if other!=person]
    scores.sort()
75    scores.reverse()
    return scores[0:n]

# Gets recommendations for a person by using a weighted average
# of every other user's rankings
80 def getRecommendations(prefs, person, similarity=sim_pearson):
    totals={}
    simSums={}
    for other in prefs:
        # don't compare me to myself
85        if other==person: continue
        sim=similarity(prefs, person, other)

        # ignore scores of zero or lower
        if sim<=0: continue
90        for item in prefs[other]:

            # only score movies I haven't seen yet
            if item not in prefs[person] or prefs[person][item]==0:
                # Similarity * Score
95                totals.setdefault(item,0)
                totals[item]+=prefs[other][item]*sim
            # Sum of similarities
```

```
        simSums.setdefault(item,0)
        simSums[item]+=sim

100
    # Create the normalized list
    rankings=[(total/simSums[item],item) for item,total in totals.items()]

    # Return the sorted list
105    rankings.sort()
    rankings.reverse()
    return rankings

def transformPrefs(prefs):
110    result={}
    for person in prefs:
        for item in prefs[person]:
            result.setdefault(item,{})

115    # Flip item and person
    result[item][person]=prefs[person][item]
    return result

120 def calculateSimilarItems(prefs,n=10):
    # Create a dictionary of items showing which other items they
    # are most similar to.
    result={}
    # Invert the preference matrix to be item-centric
125    itemPrefs=transformPrefs(prefs)
    c=0
    for item in itemPrefs:
        # Status updates for large datasets
        c+=1
130    if c%100==0: print "%d / %d" % (c,len(itemPrefs))
        # Find the most similar items to this one
        scores=topMatches(itemPrefs,item,n=n,similarity=sim_distance)
        result[item]=scores
    return result

135 def getRecommendedItems(prefs,itemMatch,user):
    userRatings=prefs[user]
    scores={}
    totalSim={}
140    # Loop over items rated by this user
    for (item,rating) in userRatings.items():

        # Loop over items similar to this one
        for (similarity,item2) in itemMatch[item]:
145
            # Ignore if this user has already rated this item
            if item2 in userRatings: continue
            # Weighted sum of rating times similarity
            scores.setdefault(item2,0)
150    scores[item2]+=similarity*rating
```

```
    # Sum of all the similarities
    totalSim.setdefault(item2,0)
    totalSim[item2]+=similarity

155 # Divide each total score by total weighting to get an average
rankings=[(score/totalSim[item],item) for item,score in scores.items( )]

    # Return the rankings from highest to lowest
rankings.sort( )
160 rankings.reverse( )
    return rankings

def calculateSimilarUser(prefs,n=5):
    # Create a dictionary of users showing which other users they
165 # are most similar to.
    result={}
    c=0
    for user in prefs:
        # Status updates for large datasets
170 c+=1
        if c%100==0: print "%d / %d" % (c,len(prefs))
        # Find the most similar items to this one
        scores=topMatches(prefs,user,n=n,similarity=sim_pearson)
        result[user]=scores
175 return result

    # Find the mean
def mean(a):
    return sum(a) / len(a)

180

def loadMovieLens():
    # Get movie titles
    movies={}
185 for line in open('u.item'):
        (id,title)=line.split('|')[0:2]
        movies[id]=title
    # Load data
    prefs={}
190 for line in open('u.data'):
        (user,movieid,rating)=line.split('\t')[0:3]
        prefs.setdefault(user,{})
        prefs[user][movies[movieid]]=float(rating)

195 # Load data gender
    gender={}
    for line in open('u.user'):
        (id,age,g) = line.split('|')[0:3]
        gender.setdefault(id,[])
200 age_g= [age,g]
        gender[id]=age_g

    ##Q1. What 5 movies have the highest average ratings?
```

```
rating_hi_avg = {}
205 for user in prefs.keys():
    for key,value in prefs[user].iteritems():
        rating_hi_avg.setdefault(key, [])
        rating_hi_avg[key].append(value)
average = {}
210 for movie in rating_hi_avg.keys():
    average[movie] = mean(rating_hi_avg[movie])

sorted_x = sorted(average.iteritems(), key=operator.itemgetter(1))
sorted_x.reverse()
215 print ("Q1 Answer")
for (key,value) in sorted_x[0:10]:
    print key, ' & ', value, '\\\\'

##Q2. What 5 movies received the most ratings?
220 lengthList = {}
for movie in rating_hi_avg.keys():
    lengthList[movie] = len(rating_hi_avg[movie])

sorted_x = sorted(lengthList.iteritems(), key=operator.itemgetter(1))
225 sorted_x.reverse()
print ("Q2 Answer")
for (key,value) in sorted_x[0:5]:
    print key, ' & ', value, '\\\\'

##Q3. What 5 movies were rated the highest on average by women?
230 rating_women = {}
for user in prefs.keys():
    if gender[user][1] == 'M':
        continue
235 for key,value in prefs[user].iteritems():
    rating_women.setdefault(key, [])
    rating_women[key].append(value)
average = {}
for movie in rating_women.keys():
240 average[movie] = mean(rating_women[movie])

sorted_x = sorted(average.iteritems(), key=operator.itemgetter(1))
sorted_x.reverse()
245 print ("Q3 Answer")
for (key,value) in sorted_x[0:11]:
    print key, ' & ', value, '\\\\'

##Q4. What 5 movies were rated the highest on average by men?
rating_men = {}
250 for user in prefs.keys():
    if gender[user][1] == 'F':
        continue
    for key,value in prefs[user].iteritems():
        rating_men.setdefault(key, [])
255 rating_men[key].append(value)
average = {}
```

```
for movie in rating_men.keys():
    average[movie] = mean(rating_men[movie])

260 sorted_x = sorted(average.iteritems(), key=operator.itemgetter(1))
sorted_x.reverse()
print ("Q4 Answer")
for (key,value) in sorted_x[0:15]:
    print key, ' & ', value, ' \\\\'

265 ##Q5. What movie received ratings most like Top Gun?
# reverse pref
item_prefs = transformPrefs(prefs)
print ("Q5 Answer")
270 print ('Most Like : ', topMatches(item_prefs, 'Top Gun (1986)')[0])
print ('Least Like : ', topMatches(item_prefs, 'Top Gun (1986)')[len(topMatches(
    item_prefs, 'Top Gun (1986)')) - 1 ])

##Q6. Which 5 raters rated the most films?
userMostRating = {}
275 for user in prefs.keys():
    userMostRating.setdefault(user, len(prefs[user]))

sorted_x = sorted(userMostRating.iteritems(), key=operator.itemgetter(1))
sorted_x.reverse()
280 print ("Q6 Answer")
for (key,value) in sorted_x[0:5]:
    print key, ' & ', value, ' \\\\'

##Q7. Which 5 raters most agreed with each other?
285 print ("Q7 Answer")
simi_users = calculateSimilarUser(prefs,1)
i=0
cumulative={}
for user in simi_users:
290     cumulative.setdefault(user,0)
    print("UserId: "), user
    num=0
    raters=[]
    print("UserValue: ")
295     print simi_users[user]
    for (key,value) in simi_users[user]:
        num+=key
        raters.append(value)
    print(num)
300     if num >= 4.0:
        i=1+i
        cumulative[user]=(num,raters)
    print cumulative[user]

305     print (i)
sorted_x = sorted(cumulative.iteritems(), key=operator.itemgetter(1))
sorted_x.reverse()
```

```

310 for (key,value) in sorted_x[0:5]:
    print key, ' & ', value, ' \\\\'

##Q8. Which 5 raters most disagreed with each other (negative correlation)

sorted_x = sorted(cumulative.iteritems(), key=operator.itemgetter(1))
315 print ("Q8 Answer")
for (key,value) in sorted_x[0:5]:
    print key, ' & ', value, ' \\\\'

##Q9. What movie was rated highest on average by men over 40?

320 rating_m_up={}
rating_m_down={}
for user in prefs.keys():
    if gender[user][0] < '40' and gender[user][1] == 'M':
325         for key,value in prefs[user].iteritems():
            rating_m_down.setdefault(key, [])
            rating_m_down[key].append(value)
        elif gender[user][0] > '40' and gender[user][1] == 'M':
            for key,value in prefs[user].iteritems():
330                 rating_m_up.setdefault(key, [])
                rating_m_up[key].append(value)
        else:
            continue
average_m_down = {}
335 for movie in rating_m_down.keys():
    average_m_down[movie] = mean(rating_m_down[movie])

sorted_md = sorted(average_m_down.iteritems(), key=operator.itemgetter(1))
sorted_md.reverse()
340 print ("Q9 A Answer")
for (key,value) in sorted_md[0:30]:
    print key, ' & ', value, ' \\\\'
average_m_up = {}
for movie in rating_m_up.keys():
345     average_m_up[movie] = mean(rating_m_up[movie])

sorted_mu = sorted(average_m_up.iteritems(), key=operator.itemgetter(1))
sorted_mu.reverse()
print ("Q9 B Answer")
350 for (key,value) in sorted_mu[0:30]:
    print key, ' & ', value, ' \\\\'

##Q10. What movie was rated highest on average by women over 40?
rating_w_up={}
355 rating_w_down={}
for user in prefs.keys():
    if gender[user][0] < '40' and gender[user][1] == 'F':
        for key,value in prefs[user].iteritems():
            rating_w_down.setdefault(key, [])
            rating_w_down[key].append(value)
360     elif gender[user][0] > '40' and gender[user][1] == 'F':

```

```
    for key,value in prefs[user].iteritems():
        rating_w_up.setdefault(key, [])
        rating_w_up[key].append(value)
365     else:
        continue

average_w_up = {}
    for movie in rating_w_up.keys():
370         average_w_up[movie] = mean(rating_w_up[movie])

sorted_wu = sorted(average_w_up.iteritems(), key=operator.itemgetter(1))
sorted_wu.reverse()

375     print ("Q10 A Answer")
    for (key,value) in sorted_wu[0:30]:
        print key, ' & ', value, ' \\\\'

average_w_down = {}
380     for movie in rating_w_down.keys():
        average_w_down[movie] = mean(rating_w_down[movie])

sorted_wd = sorted(average_w_down.iteritems(), key=operator.itemgetter(1))
sorted_wd.reverse()

385     print ("Q10 B Answer")
    for (key,value) in sorted_wd[0:30]:
        print key, ' & ', value, ' \\\\'

390     return prefs

loadMovieLens();
```



## References

- [1] GroupLens. MovieLens. <http://grouplens.org/datasets/movielens/>